


# Gradient Coding

Rashish Tandon\*, Qi Lei\*, Alexandros Dimakis\*, Nikos Karampatziakis†

\* The University of Texas at Austin, † Microsoft Research



Neural Information  
Processing Systems  
Foundation

## Overview

We propose a novel coding theoretic framework for mitigating stragglers in distributed learning. We show how to achieve tolerance to **failures/stragglers** for **synchronous Gradient-based methods** by:

- Replicating data blocks across machines
- Coding over the gradients transmitted

## Introduction


Given a dataset  $D = \{(x_1, y_1), \dots, (x_d, y_d)\}$ , many ML tasks have the form:

$$\beta^* = \underset{\beta \in \mathbb{R}^p}{\operatorname{argmin}} \sum_{i=1}^d \ell(\beta; x_i, y_i) \quad (1)$$

Gradient-based methods solve this via the update step:

$$\beta^{(t+1)} = h(\beta^{(t)}, g^{(t)}) \quad (2)$$

- $g := \sum_{i=1}^d \nabla \ell(\beta^{(t)}; x_i, y_i)$ , the **gradient**
- If  $d$  is large, its computation can be distributed
- Several methods e.g. gradient descent, accelerated gradient descent, conditional gradient, LBFGS, bundle methods etc. fit in this framework
- **Notation:**  $d$  samples,  $n$  workers denoted as  $\{W_1, \dots, W_n\}$ ,  $s$  stragglers (with  $s < n$ ), Data partitions denoted as  $\{D_1, \dots, D_n\}$



## General Setup

We seek matrices  $A \in \mathbb{R}^{f \times n}$ ,  $B \in \mathbb{R}^{n \times n}$ :

$$AB = \mathbf{1}_{f \times n} \quad (3)$$

- $f$  is number of combinations of non-straggler workers
- $\mathbf{1}_{f \times n}$  is the all 1s matrix of size  $f \times n$
- $B_i$  ( $i^{\text{th}}$  row of  $B$ ) is *linear combination* over **partial gradients** computed by  $i^{\text{th}}$  worker
- $A_i$  ( $i^{\text{th}}$  row of  $A$ ) is *linear combination* over **workers** when  $\operatorname{supp}(A_i)$  survive
- In the above example:
  - $A = \begin{pmatrix} 0 & 1 & 2 \\ 1 & 0 & 1 \\ 2 & -1 & 0 \end{pmatrix}$ ,  $B = \begin{pmatrix} 1/2 & 1 & 0 \\ 0 & 1 & -1 \\ 1/2 & 0 & 1 \end{pmatrix}$

## Full Stragglers

- Stragglers are allowed to be **arbitrarily slow**, to the extent of complete failure
- Rows of  $A$  have all subsets over  $[n]$  of size  $(n - s)$  as supports ( $f = \binom{n}{n-s}$ )

We require that for any set  $I \subseteq [n]$  s.t.  $|I| = n - s$ ,  $B$  satisfies:

$$\mathbf{1}_{1 \times n} \in \operatorname{span} \{B_i \mid i \in I\} \quad (4)$$

- The above condition is necessary and sufficient to get  $(A, B)$  s.t.  $AB = \mathbf{1}$

**Theorem (Lower Bound on B's density):** If all rows of  $B$  have the same no. of non-zeros, then:

$$\|B_i\|_0 \geq (s + 1) \text{ for all } i \in [n]$$

## Fractional Repetition Scheme:

- Applicable when  $n$  is a multiple of  $(s + 1)$
- Partition data equally among  $\frac{n}{(s+1)}$  workers
- Create  $(s + 1)$  replicas using other workers
- Every worker sends sum of its partial gradients

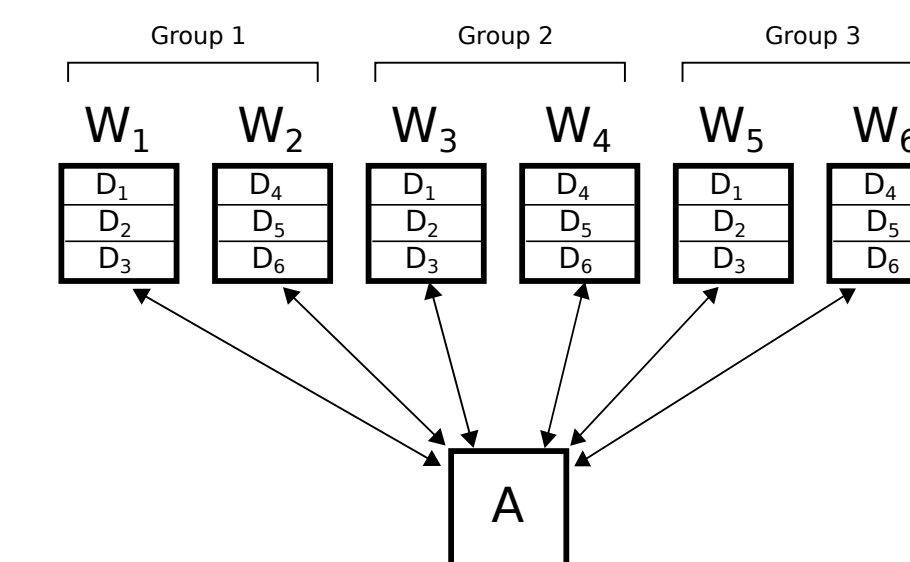


Figure: Fractional Repetition Scheme for  $n = 6, s = 2$

## Cyclic Repetition Scheme: $B$ has the support structure

$$\operatorname{supp}(B) = \begin{bmatrix} \overbrace{\star \star \cdots \star}^{s+1} & \star & 0 & 0 & \cdots & 0 & 0 \\ 0 & \star & \star & \cdots & \star & \star & 0 & \cdots & 0 & 0 \\ \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \ddots & \ddots & \vdots & \vdots \\ 0 & 0 & \cdots & 0 & 0 & \star & \star & \cdots & \star & \star \\ \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \ddots & \ddots & \vdots & \vdots \\ \star & \cdots & \star & \star & 0 & 0 & \cdots & 0 & 0 & \star \end{bmatrix}_{n \times n} \quad (5)$$

- Pick rows of  $B$  from a random subspace  $S$  (described as null-space of an  $s \times n$  MDS matrix) and also containing  $\mathbf{1}_{n \times 1}$
- Always possible for any  $n$  and  $s$

## Partial Stragglers

- Stragglers are allowed to be **at most  $\alpha$ -times slower** than any non-straggler
- Key idea is to couple the naive scheme with the coding scheme
- Data is split into a *naive* component and *coded* component
- When a straggler finishes processing its *naive* components, a non-straggler finishes its *naive+coded* components

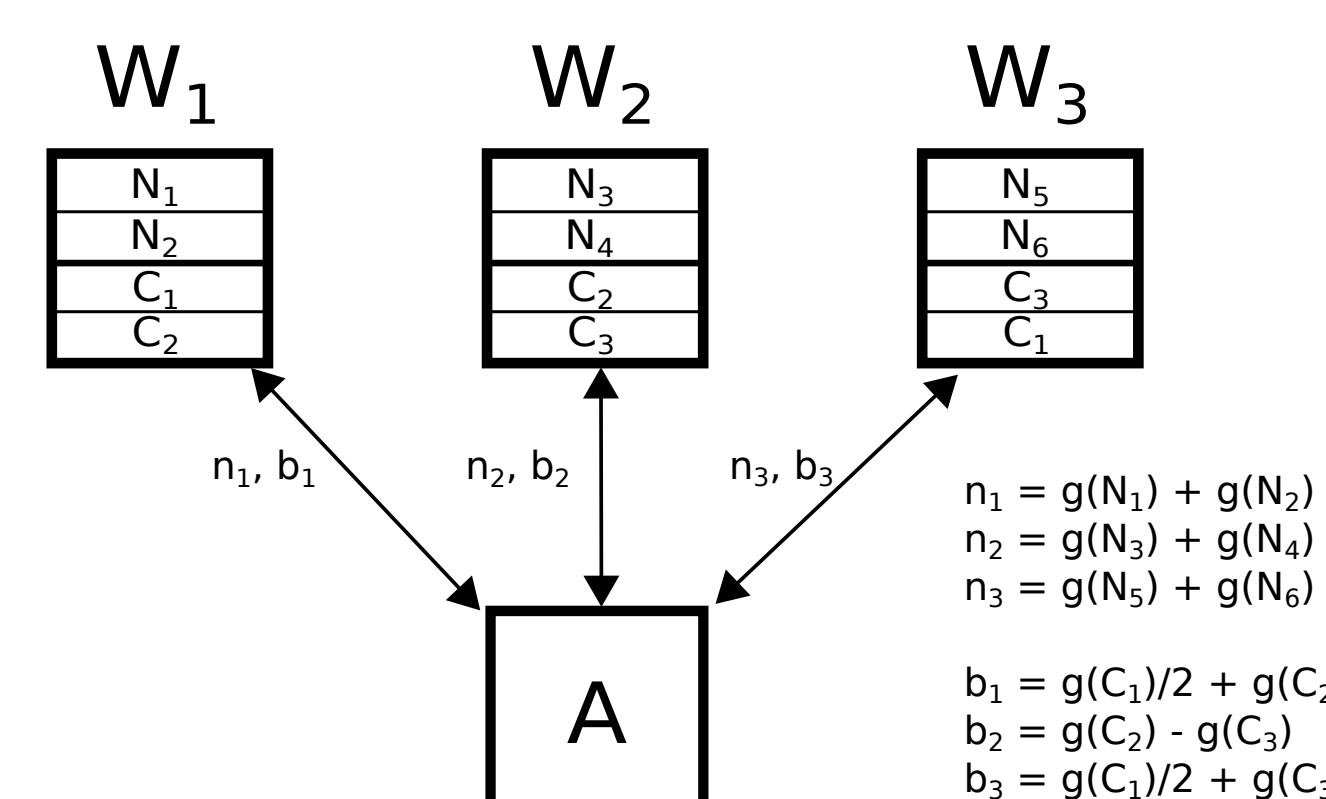
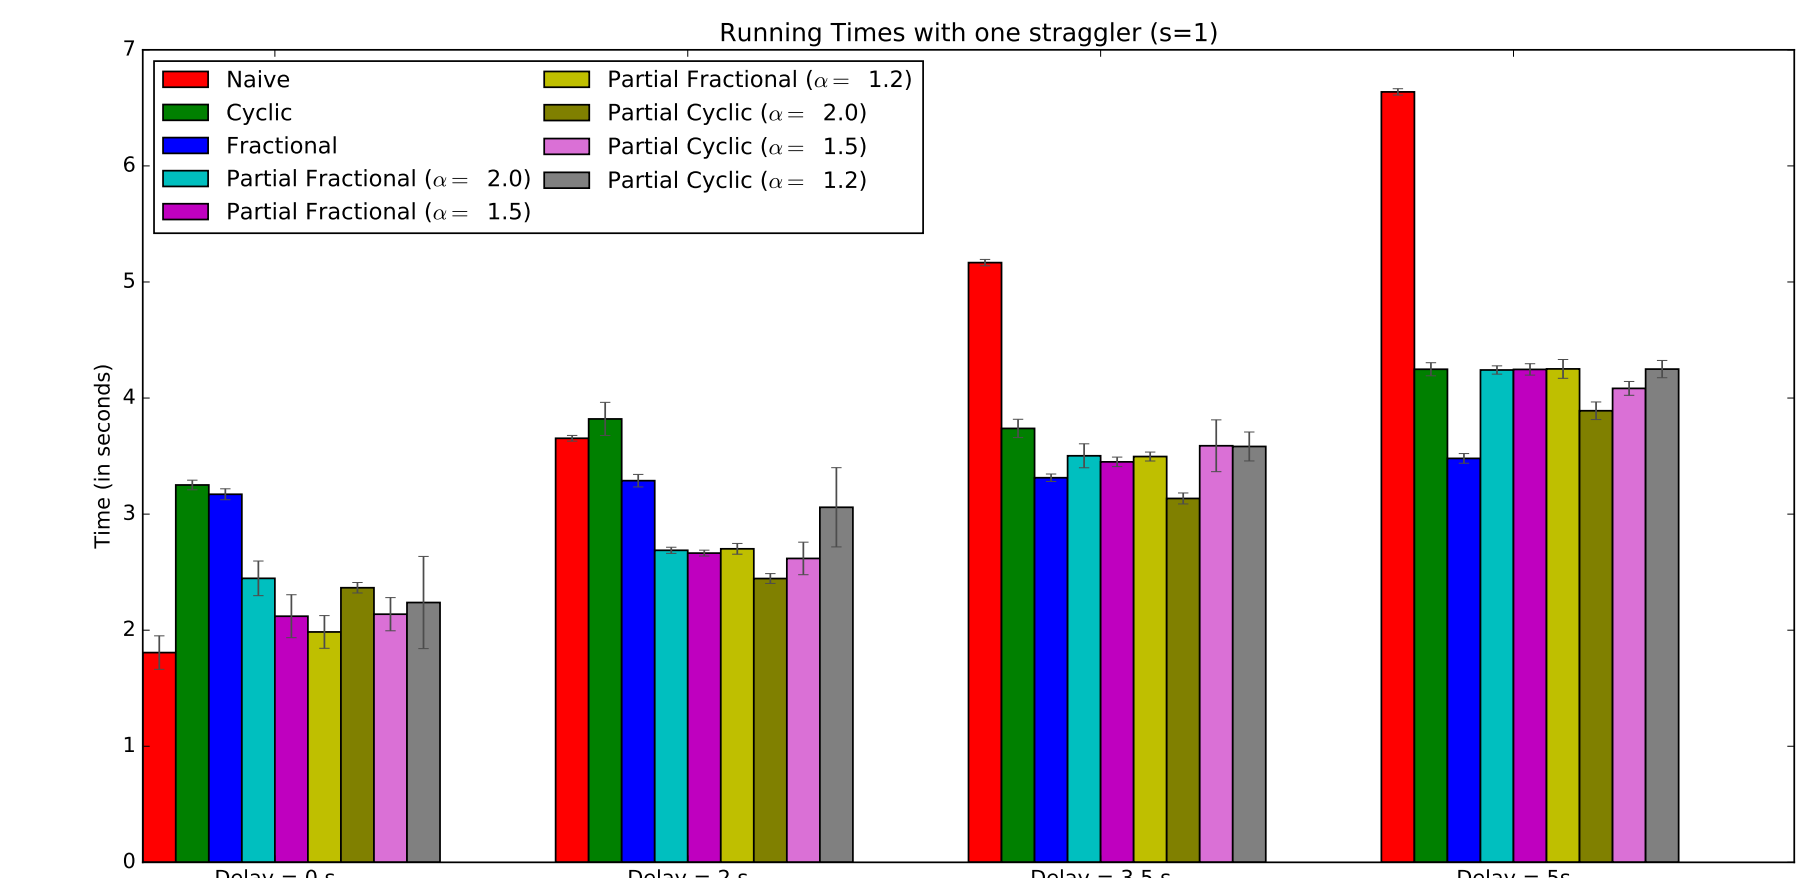


Figure: Scheme for Partial Stragglers,  $n = 3, s = 1, \alpha = 2$ .  $g(\cdot)$  represents the partial gradient. Each worker gets  $4/9$  fraction of the data (as opposed to a  $2/3$  fraction in the earlier scheme).

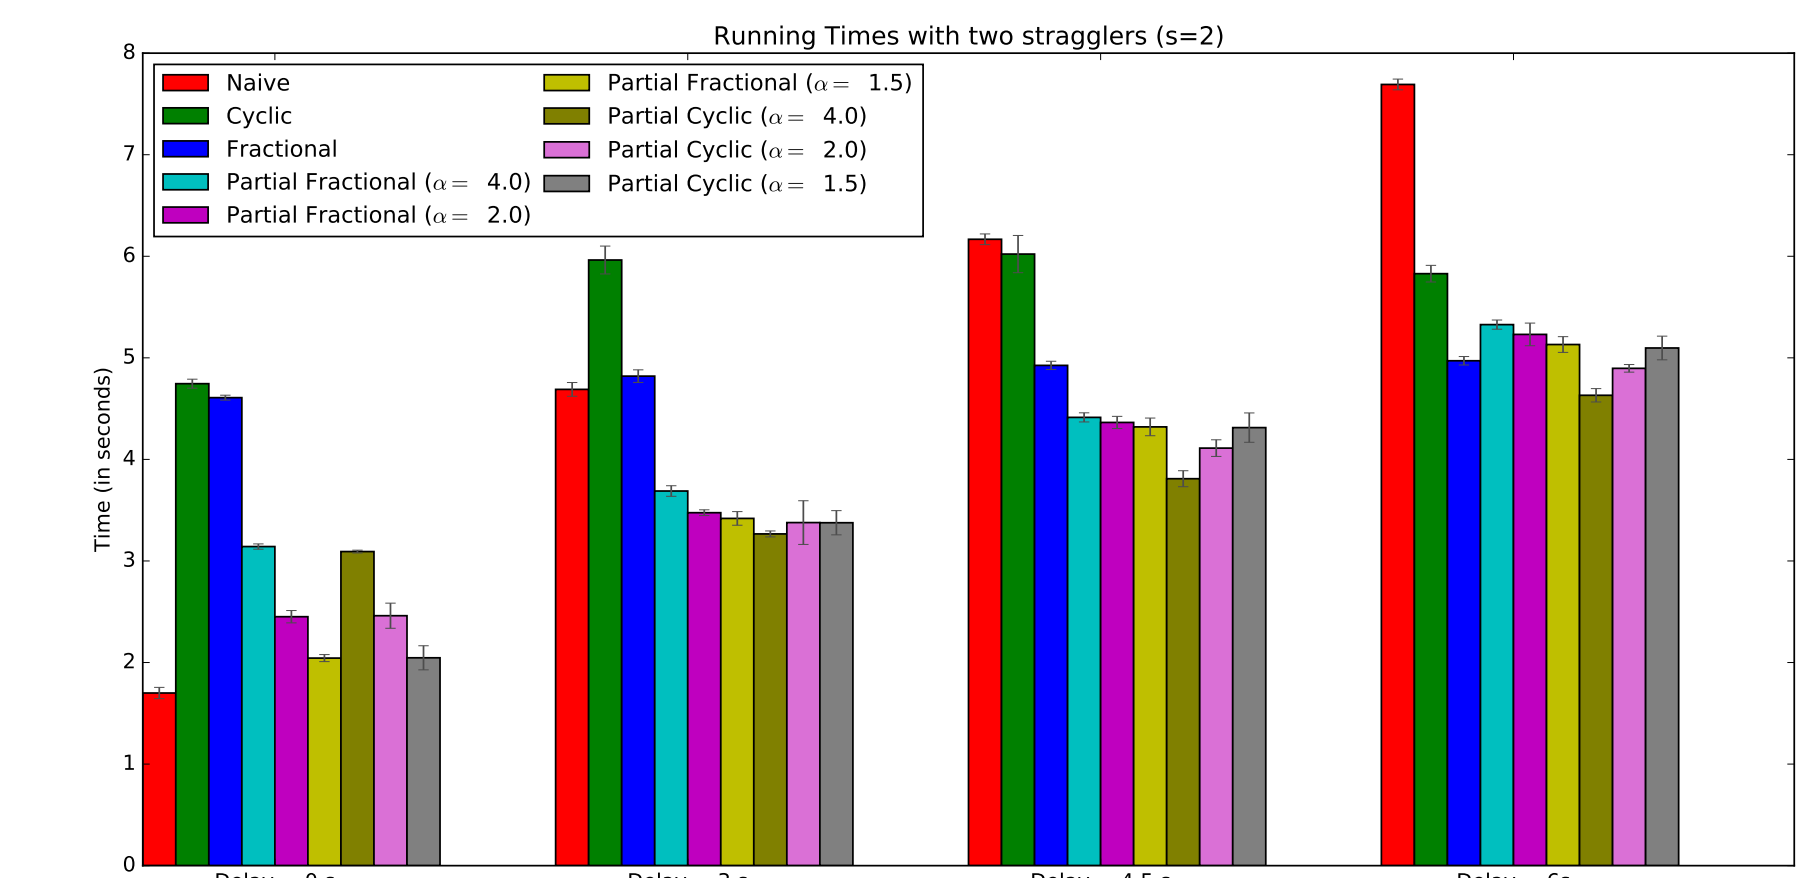
## Experiments

### Artificial Dataset:

- Generate a dataset with  $d = 554400$  samples,  $D = \{(x_1, y_1), \dots, (x_d, y_d)\}$
- $x \sim \frac{1}{2} \times \mathcal{N}(\mu_1, I) + \frac{1}{2} \times \mathcal{N}(\mu_2, I)$  and  $y \sim \operatorname{Ber}(p)$ ,  $p = 1 / (\exp(2x^T \beta^*) + 1)$
- Model dimension  $p = 100$ , and  $\mu_1, \mu_2, \beta^*$  chosen randomly



(a) One straggler

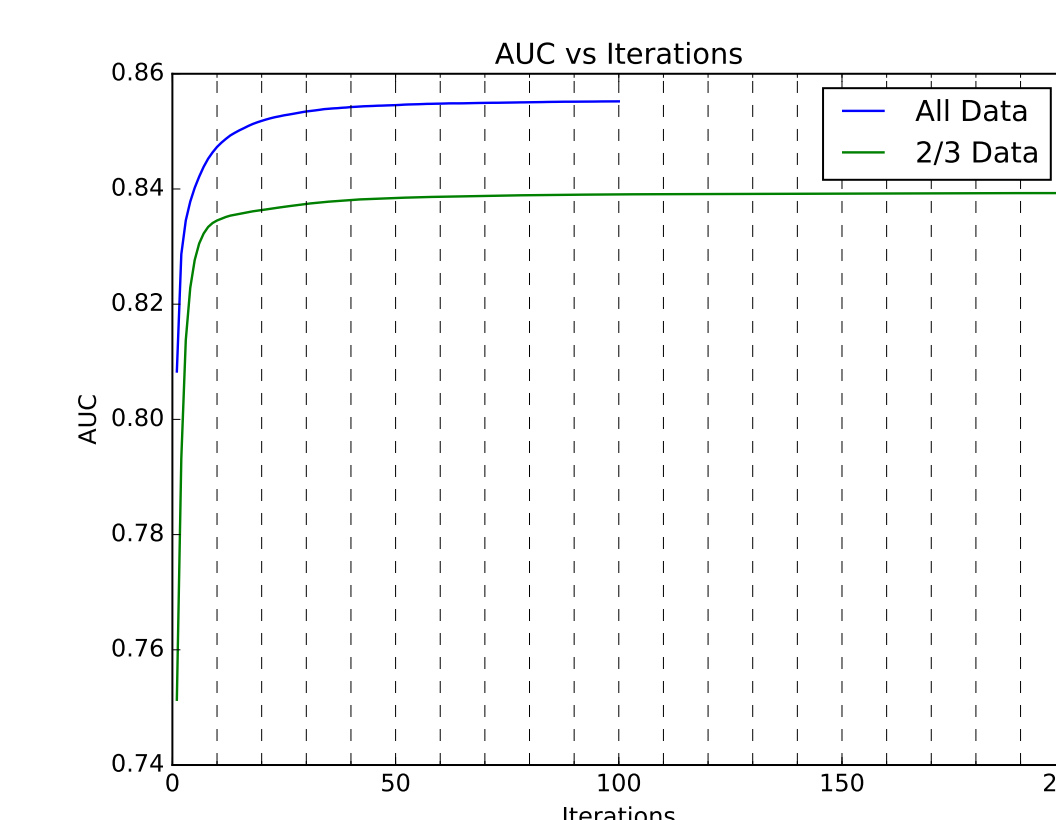


(b) Two stragglers


Figure: Empirical running time on EC2 with  $n = 12$  small machines, and one or two stragglers. Other machines run at normal speed while stragglers are artificially delayed.  $\alpha$  is slowdown rate for partial schemes. We note that the partial straggler schemes for  $\alpha = 1.2$  needs to only replicate approximately 10% of the blocks.

### Real Dataset:

- Train a logistic regression model on a real binary classification dataset
- Generalization error comparison with *ignoring stragglers* approach i.e. data is distributed equally and we only wait for  $(n - s)$  machines to finish



(a)  $n = 3, s = 1$ : Missing 1/3 of the data



(b)  $n = 5, s = 1$ : Missing 1/5 of the data

Figure: Generalization Error comparison of our scheme vs. *ignoring stragglers*. The straggling worker was fixed to be the same machine across iterations.